



Understanding Microservices Architecture

Maaïke van Putten
Software developer & trainer
www.biltup.com



Y

Session 1

Introduction to
microservices

d

Session 2

Comparing
microservices and
monolithic
architecture

d

Session 3

Building and
containerizing
microservices

E

Session 4

Managing data in
microservices



Managing data in microservices



Learning objectives

u

Understand how to **handle data persistence** in Docker.

d

Implement a simple **database service** (like PostgreSQL) as a separate container.

F

Connect the **database** service to a **Python microservice** for CRUD operations.

Bounded context pattern

b

Splitting up a domain into subdomains

R

Each entity has an unambiguous meaning within the subdomain

W

Different subdomains can have different properties for same named entities



Mappers for entities for communication between subdomains




Bounded contexts in microservices are typically implemented as individual microservices

d

They shouldn't share a common database



Data model



Different names for entities across subdomains to best match the usage of the entity



Unique storage



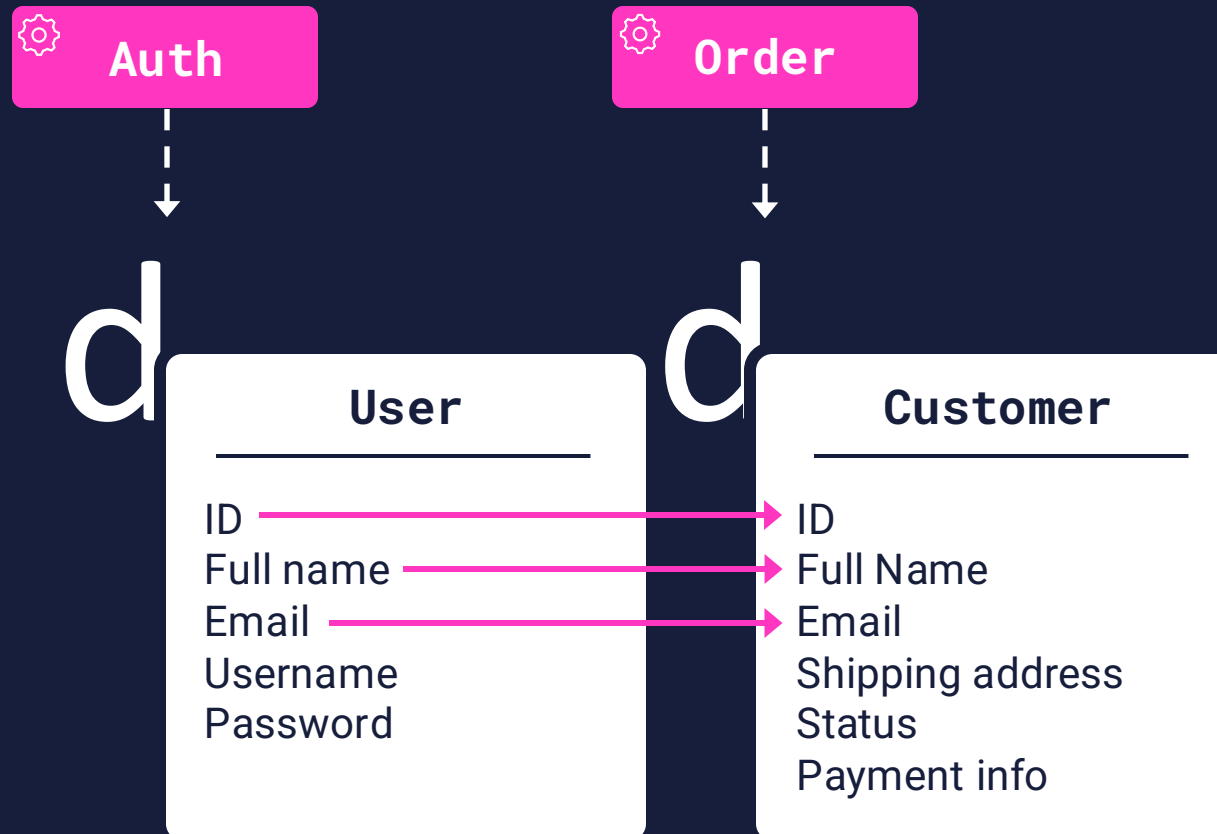
Consistency and syncing data is a complication



Complex queries that are sourced from different databases across microservices are harder

Example bounded context

User and customer, different interesting properties for each microservice



Bounded context and microservice

- Splitting up the landscape into bounded contexts is a great way to start
- Often a bounded context is mapped to a single microservice, but this doesn't have to be the case
- It also depends on how independent the contexts are in terms of business logic and data ownership
- Also take non-function requirements into account:
 - Be able to independently scale
 - Release frequency and need for changes
 - Separate team for microservice
 - Best choice for the tools (tech stack)
 - Data ownership



Data design considerations

E

Microservices should manage their own data

d

They should not share a database

h

APIs can be used to get data from the database of another microservice

f

BUT: API calls are costly and they shouldn't be constantly happening

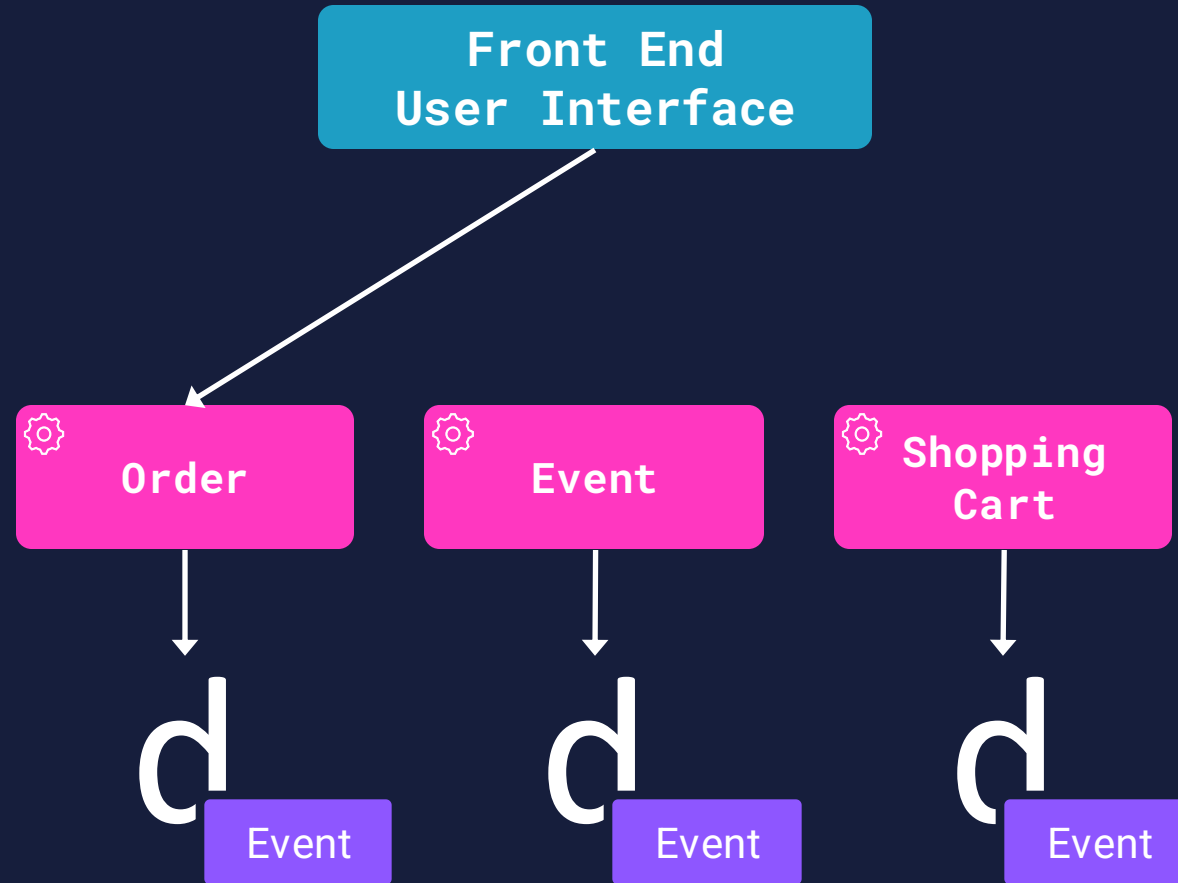
l

Async calls and service bus for delivering the messages can help (also with resilience if the other service is down for a short time)

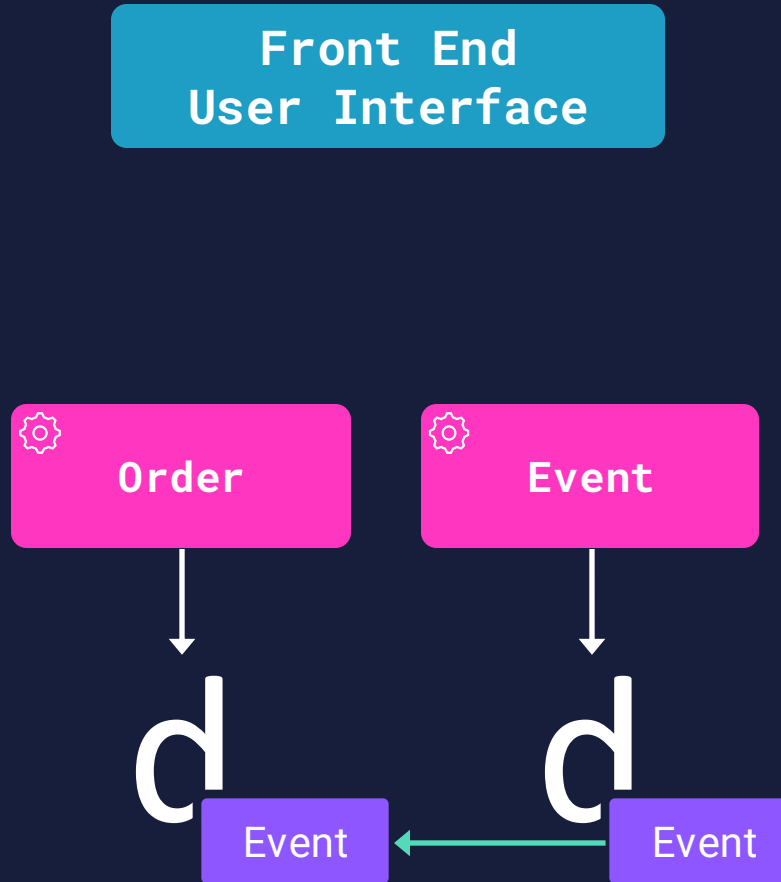
p

If two microservice need realtime data from each other and delay cannot be tolerated, and/or have a lot of data that they share in general >> one microservice instead of two

Duplicate data




Duplicate data



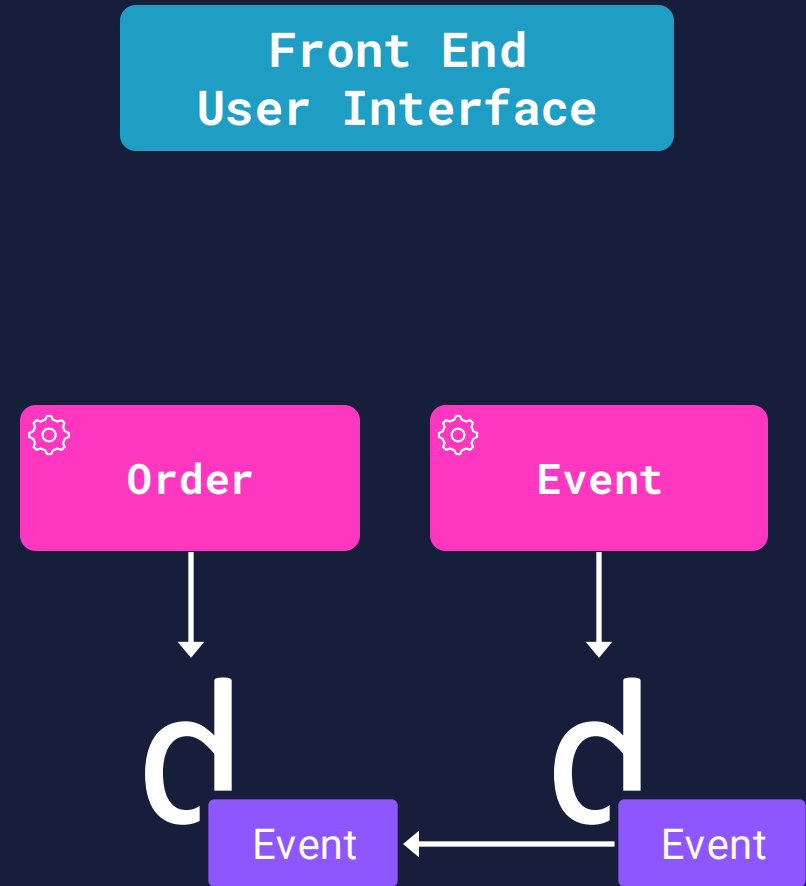
Polygot persistence

- Type of data storage might differ per microservice (for example mixing noSQL and SQL)
- This mixing of data storage techniques is called “polygot persistence”




Eventual consistency

- Related data across microservices has to be synced
- Eventual consistency is the name for the most common strategy
- Changes are synced but this takes a little bit (no blocking of the microservices during the update)



Implementing eventual consistency

- Event driven communication with publishers and subscribers
- Async
- System needs to tolerate that not all data is immediately consistent




Queries with distributed data

Three common options:

- API gateway
- Materialized views
- Cold data in central databases




API gateway (another microservice itself)



Materialized view

- View of the data in one microservice
- Could give consistency issues



Central reporting database



Cold data (data that is not being changed but only read) can be stored in a central database, such as a data warehouse



This is used for reporting and not for changing data



Again, consistency issues and eventual consistency approach

Data consistency in distributed systems



Failures can always occur, distributed architecture has multiple pain points



It's possible that calls between services fail, typically circuit breaker pattern used for this



Service bus also helps to mitigate the issues



Integration events for eventual consistency

What we'll be doing:



d

Run a database in a container

p

Persist the data in a volume

U

Interact with the containerized database from our microservices

Exercise

Adding data persistence to a database service



Exercise

Connect microservices from last time to the database





Questions?

maaikejvp@gmail.com



www.biltup.com